

Diffusion Cartograms in the Wolfram Language

A self-contained pure-Wolfram implementation of the Gastner–Newman density-equalising map algorithm.

Marco Thiel, April 2026 — github.com/mthiel74/Contiguous-Cartograms

Abstract

Cartograms are maps in which the area of each region is deformed to be proportional to a statistical quantity — population, GDP, number of COVID cases, votes cast, and so on — while the topology (who is next to whom) is preserved. Gastner and Newman (2004) showed how to produce such maps by interpreting the quantity as a density and letting it diffuse to uniformity; every point of the map is advected by the resulting velocity field $\mathbf{v} = -\nabla \rho / \rho$. This notebook presents a compact, pure-Wolfram-Language port of that algorithm. It ships as a paclet-style package `CartogramWL` exposing both low-level primitives (`DiffusionSolver`, `Cartogram`, `CartogramTransform`) and a one-liner `WorldCartogram[metric]` that pulls country polygons and statistical values from the built-in `CountryData` and returns a ready-to-publish figure. It supports any `ColorData` gradient, a `GeoBackground` → "Satellite" mode that textures each country with deformed satellite imagery, user-supplied `Association` input, and visualisation of the underlying space deformation via a distortion grid.




Figure 1. Left: countries on the geographic lon/lat grid, each textured with satellite imagery. Right: the same countries with area proportional to GDP, the satellite imagery deformed along with the geometry. Ocean area is preserved.

1. What is a cartogram?

A geographic map of the world devotes enormous screen area to Russia, Canada, and the Sahara, and a few pixels each to Bangladesh, South Korea, or the Netherlands. If the question being asked is "where do people live?" or "where is the economy?", this is the opposite of what we want. Density-equalising cartograms solve this by reshaping each region so that its area tells the story of a variable of interest.

Gastner & Newman's elegant insight (PNAS 2004) is that a map in which density $\rho(\mathbf{r})$ is uniform is exactly what we want, and diffusion is the natural process that takes any non-uniform density to a uniform one. We therefore simulate the heat equation

$$\frac{\partial \rho(\mathbf{r}, t)}{\partial t} = \nabla^2 \rho(\mathbf{r}, t), \quad \nabla \rho(\mathbf{r}, t) \cdot \mathbf{n} = 0 \quad \text{on } \partial \Omega$$

on a rectangular domain with Neumann boundary conditions (no flux through the box edge, so mass is conserved), and flow each geographic point through the instantaneous velocity

$$\frac{d\mathbf{r}}{dt} = \mathbf{v}(\mathbf{r}, t) = -\frac{\nabla \rho(\mathbf{r}, t)}{\rho(\mathbf{r}, t)}.$$

As $t \rightarrow \infty$ the density relaxes to its spatial mean and the advected points settle into their cartogram positions. Regions of high initial density draw space inwards and therefore expand; regions of low density are pushed outwards.

2. The algorithm, step by step

2.1 Rasterise polygons to a density grid

Given a collection of country polygons and a per-country metric M_c (population, GDP, ...), we set the density inside polygon c to

$$\rho_0(\mathbf{r}) = \frac{M_c}{A_c} \quad \text{for } \mathbf{r} \in c,$$

where A_c is the geographic area of c . Integrated over c we recover the metric M_c . Ocean cells are filled with the mean land density so the spatial mean is uniform across the domain and oceans retain their size; only the relative sizes of countries change.

The rasterisation uses a **RegionMember**-based subpixel test for anti-aliasing on thin/small countries.

2.2 Solve the diffusion equation

The Laplacian on a rectangle with Neumann boundary conditions is diagonalised by the cosine basis. Wolfram's built-in `FourierDCT[... , 2]` and `FourierDCT[... , 3]` implement the DCT-II and its inverse in any number of dimensions with an orthonormal convention, so the full 2D spectrum is one call each. Expanding

$$\rho(\mathbf{r}, t) = \sum_{m, n} \hat{\rho}_{mn} e^{-\lambda_{mn} t} \cos\left(\frac{\pi m x}{L_x}\right) \cos\left(\frac{\pi n y}{L_y}\right)$$

with eigenvalues $\lambda_{mn} = \pi^2 \left(\frac{m^2}{L_x^2} + \frac{n^2}{L_y^2} \right)$, the density at any time t is evaluated in $\mathcal{O}(N \log N)$ from the cached cosine coefficients $\hat{\rho}_{mn}$.

2.3 Advect points through the velocity field

The velocity on the grid is

$$\mathbf{v}(\mathbf{r}, t) = -\frac{\nabla \rho(\mathbf{r}, t)}{\rho(\mathbf{r}, t)},$$

computed by central differences on the current diffused grid and interpolated bilinearly at any query point. Each geographic point is integrated from $t = 0$ to $t^{\star}$ using an explicit Runge–Kutta scheme (`NDSolveValue[... , Method -> "ExplicitRungeKutta"]`) with a 5th-order stepper, where $t^{\star}$ is chosen from the decay constant of the slowest non-constant Fourier mode.

2.4 Pre-smoothing

Because polygon-rasterisation produces step functions, the initial velocity field is very stiff. We pre-smooth the density with a small Gaussian filter ($\sigma \sim 1$ cell), which is mathematically equivalent to starting the diffusion at $t_0 = \sigma^2 \Delta x \Delta y / 2$. A small background floor (`MeanFloor`) is added to avoid division by zero inside zero-density cells.

3. The CartogramWL package

3.1 Loading

Clone the repository from GitHub and load the package directly from the `CartogramWL/` directory. The package has zero external dependencies.

```
Get["/path/to/wolfram/CartogramWL/CartogramWL.wl"]
```

All public symbols are then available in the `CartogramWL`` context:

```
Names["CartogramWL`*"]
```

3.2 Low-level primitives

The algorithm is exposed as a handful of composable functions that all operate on a density array `rho[[i, j]]` with row index $i \rightarrow y$, column index $j \rightarrow x$ (NumPy convention) and a bounding box `{xmin, ymin, xmax, ymax}`:

```
solver = DiffusionSolver[rho, {xmin, ymin, xmax, ymax}];
rhoAtT = DensityAt[solver, t];
newPts = AdvectPoints[solver, {{x1, y1}, {x2, y2}, ...}];
```

One level up, `Cartogram` bundles the pre-smoothing, mean floor, and ocean-preservation settings; `CartogramRun` precomputes the deformed cell-centre grid; `CartogramTransform` maps arbitrary points through the deformation; `CartogramTransformPolygon` additionally densifies long edges before the transform.

```
cart = Cartogram[rho, bbox,
  "MeanFloor" -> 0.02, "BlurSigma" -> 1.5,
  "SeaDensity" -> "auto"];
cart = CartogramRun[cart, "Tol" -> 2.*^-3];
newPoly = CartogramTransformPolygon[cart, polyRing];
```

3.3 The one-liner: WorldCartogram

Most end users want "given a metric, draw the world". `WorldCartogram` wraps the whole pipeline, reads polygons and values from `CountryData`, and returns a two-panel `GraphicsRow` (geographic + cartogram).


```
WorldCartogram["Population"]      (* or "GDP", "GDPPerCapita", ... *)
WorldCartogram["GDP", "ColorFunction" → ColorData["TemperatureMap"]]
WorldCartogram["GDP", "Background" → "Satellite"]
WorldCartogram[{"my metric", <|"Germany" → 83., ...|>}]
```

A complete option list is available via `?WorldCartogram`. The most useful options are `"ColorFunction"` (any `ColorData[...]` gradient or a raw function on $\$[0,1]\$$), `"Background"` (`"Flat"` or `"Satellite"`), `"MissingCountries"` (`"Hide"` or `"ShowGrey"`), and `"BoundingBox"` (useful for continent-scale maps).


4. Built-in metrics

4.1 Population

The canonical cartogram: India, China, and sub-Saharan Africa bloom outward; Russia, Canada, and Australia collapse.




4.2 GDP




4.3 GDP per capita

GDP per capita is a per-area quantity, so the cartogram rewards wealth density rather than raw wealth. Luxembourg, Switzerland, the Gulf states and Scandinavia visibly bulge; Africa and parts of South Asia contract.




4.4 Population density



5. Colour schemes & satellite background


Pass any `ColorData` gradient as "ColorFunction":

```
WorldCartogram["PopulationDensity",
  "ColorFunction" → ColorData["TemperatureMap"]]
```



Set "Background" → "Satellite" and each country is textured with the corresponding patch of Wolfram's `GeoBackground` → "Satellite" imagery. On the cartogram side the imagery is deformed along with the country geometry, giving an effect similar to a shaded-relief globe being inflated or deflated locally.

```
WorldCartogram["GDP", "Background" → "Satellite"]
```



6. Visualising the distortion itself

A clean way to see the geometry of the transformation is to overlay a regular 15° lat/lon grid on both panels. On the left the grid is undistorted; on the right the $\$325$ grid intersections are pushed into the cartogram by the same velocity field that reshaped the countries. The grid therefore encodes the local stretching/compression factor at every point of the domain.




Figure 2. Left: geographic frame, regular 15° grid in red, countries grey. Right: the GDP cartogram, the same grid now warped by the same diffusion flow. Grid cells over high-wealth regions shrink (fewer dollars per square millimetre); cells over oceans and sparsely populated continents stretch.

`CartogramDistortionGrid[cart]` returns the `Line` primitives for the warped grid; wrap in `Graphics` to draw, overlay on a `WorldCartogram` result to annotate a cartogram with its own distortion field.

7. Bringing your own data

Any `Association` of country-identifier $\rightarrow$ value works in place of a metric string. Keys may be `CountryData` identifiers (e.g. `"UnitedStates"`) or display names (`"United States"`).

```
g20 = <|"UnitedStates" -> 332., "China" -> 1425.,
      "India" -> 1428., "Japan" -> 125., "Germany" -> 83.,
      "UnitedKingdom" -> 67., "France" -> 67.5, ...|>;
```

```
WorldCartogram[{"G20 population (millions)", g20},
  "MissingCountries" -> "ShowGrey"]
```




Figure 3. A cartogram driven by a hand-curated `Association` of 19 G20-ish economies; every other country appears in grey and does not contribute to the density field, but still visibly deforms because diffusion is global.

8. Implementation notes

- Only the diffusion is done in Fourier space; advection uses an RK45 stepper over the reconstructed velocity. Modes below a relative tolerance (10^{-3} by default) are considered extinct, giving $t^{\star}$.
- Heavy-tailed metrics (Monaco on population density, Luxembourg on GDP per capita) produce

enormous per-cell densities that can destabilise the advection. `WorldCartogram` clamps values at a configurable fraction of the mean and caps the rasterised grid at a multiple of the positive-cell mean.

- Long country edges are densified before being transformed, matching `shapely.segmentize` on the Python side of this project.
- The algorithm is grid-convention-compatible with the Python implementation so the two are drop-in interchangeable for debugging/cross-validation.

9. References

- M. T. Gastner & M. E. J. Newman, "Diffusion-based method for producing density-equalizing maps", *PNAS* **101**(20), 7499–7504 (2004). doi:10.1073/pnas.0400280101.
- Source code and the Python twin:
github.com/mthiel74/Contiguous-Cartograms.
- Relies on Wolfram's `CountryData`, `GeoGraphics`, `FourierDCT`, and `NDSolveValue`.